



# *Overview of Operating Systems*

---

## CS201 Lecture 1

Jason Hibbeler

University of Vermont

Fall 2019



# Topic Outline

---

- Complete and total overview an operating system
- Processes
- Memory and protection
- Storage



# Basic Role of an OS

---

Basic mission:

- manage resources for users of the system.

What are the resources?



# Basic Role of an OS

---

Basic mission:

- manage resources for users of the system

What are the resources?

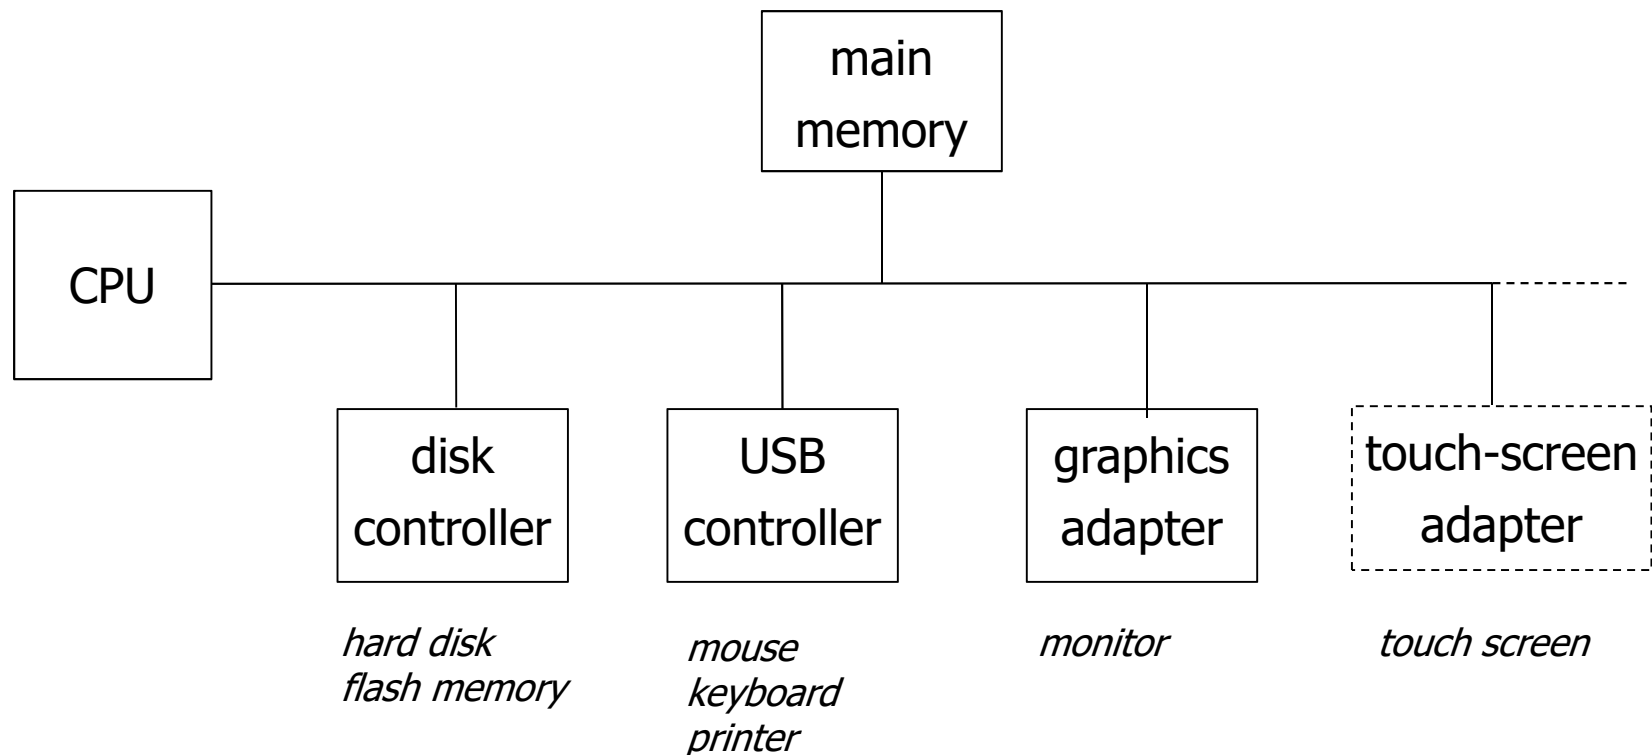
- CPU, memory, storage, I/O devices

Who wants/needs to use these resources?

# The System

What is a system?

- PC, Mac, phone, iPhone, supercomputer, mainframe, RPi, etc.



# Operating System

What is an operating system?

- it's the one program that is always running on the system
- the core part is called "the kernel"
- controls the hardware
- facilitates the execution of user programs
- also has other auxiliary programs to help out



*the Colonel*



# System Initialization

---

Turn on the power, and...

- the system loads the “bootstrap image” from ROM (read-only memory)
- this bootstrap image is also called the firmware

The bootstrap image then loads the OS

- and starts various auxiliary processes (“daemons”, for Unix and macOS, “services” for Windows)



# The OS in Operation

---

Wait for something to happen:

- wait for an event
- then handle the event
- an event is usually signaled by an interrupt—kind of a low-level “hello...I need attention!”





# Interrupt Handling

---

When something happens, take care of it

- the CPU stops what it was doing and transfers to a special location: the interrupt service code for that interrupt
- this operation must be fast!
- it's implemented in hardware by loading an interrupt vector: an array of addresses indexed by unique device number

# Interrupt Handling

Must be fast!



*interrupt comes in on a special dedicated line*

*processing ... processing ... processing ... processing*

*processing ... processing ... processing ... processing*



interrupt #1	address A1
interrupt #2	address A2
interrupt #3	address A3

■  
■  
■

CPU jumps to interrupt-handler routine

- looks up address of code to handle that specific interrupt
- saves state associated with current task
- jumps to the specified address
- runs that code
- returns and restores saved state



# Storage

---

## Main memory

- random-access memory (RAM)
- sometimes it's SRAM, sometimes DRAM
- key attributes are that it's limited in size
- and it's volatile (turn off power  $\Rightarrow$  lose contents of memory)

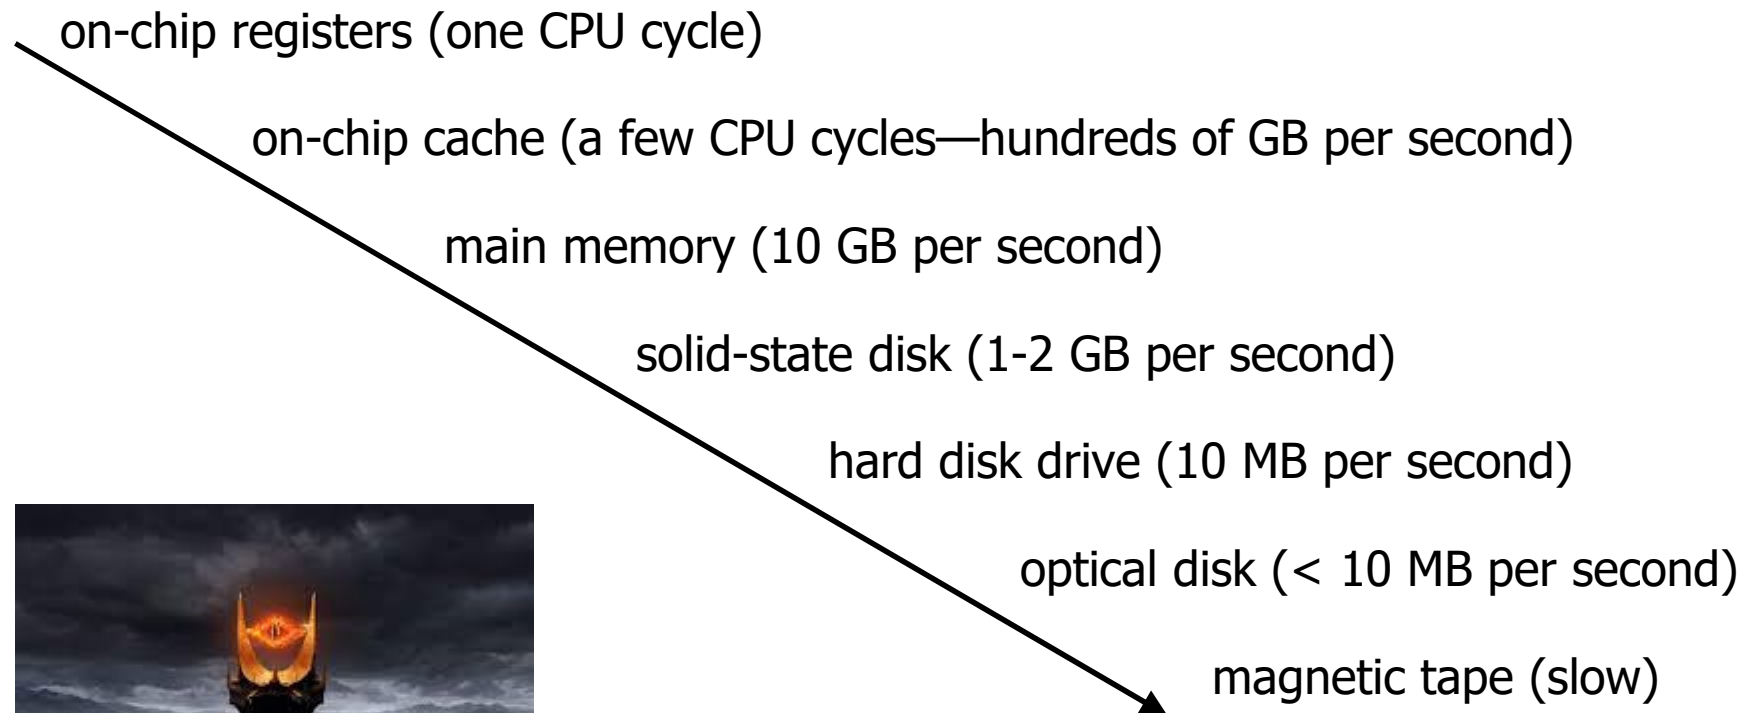
Memory holds program instructions and program data

Also, there is read-only memory (ROM, EEPROM)

By the way, what is a byte? What is a "megabyte"?  
What is a "gigabyte"?

# Memory Hierarchy

How long is a  
CPU cycle?



*the CPU*



# I/O Processing

---

Each hardware device has a device controller that is responsible for the actual interface to the hardware

OS has a device driver that talks to the device controller

Device driver loads registers in the device controller with instructions (“read a keystroke”)

- device controller then signals to the OS, with an interrupt, that’s it’s done




# Special Case: DMA

---

Direct memory access: the device controller can transfer a block of data directly to main memory

- without the intervention of the CPU
- why is this useful?



# Computer Systems Architecture

---

In the beginning, there was the single-processor system

- one CPU, doing one thing at a time
- there are usually auxiliary specialized processors as well (disk controller, keyboard controller, etc.)

Nowadays: multiprocessor (or multicore) systems are the norm

- note that “core” and “processor” are not synonymous
- multicore systems are truly the norm now

# Apple A11 Bionic

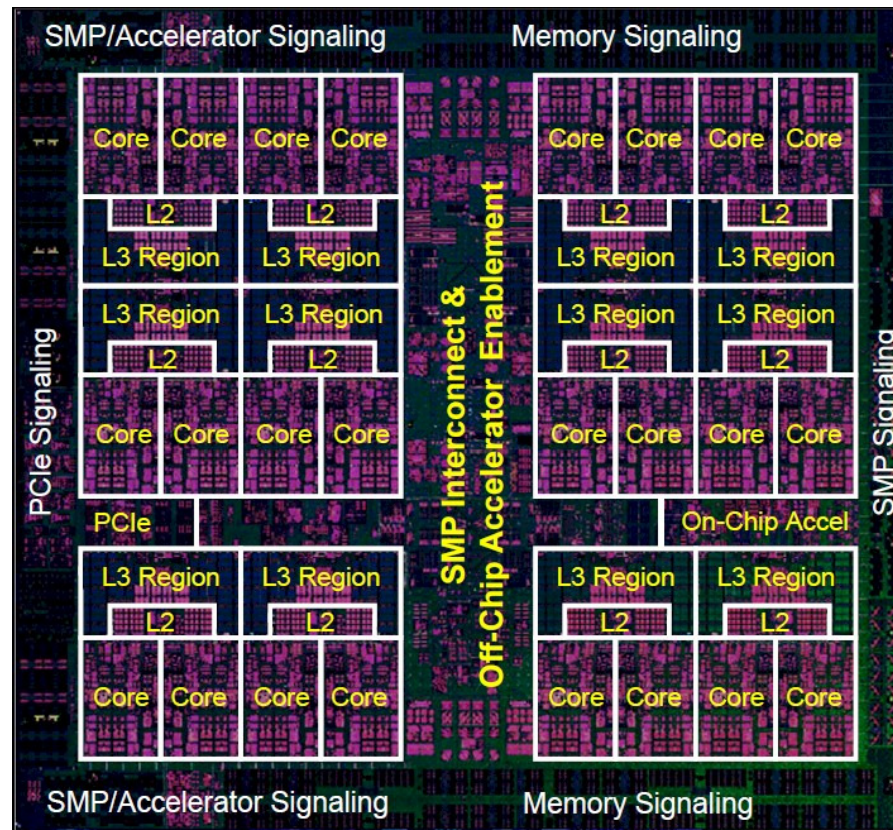


The A11 features an Apple-designed 64-bit ARMv8-A six-core CPU, with two high-performance cores at 2.39 GHz, called **Monsoon**, and four energy-efficient cores, called **Mistral**. The A11 uses a new second-generation performance controller, which permits the A11 to use all six cores simultaneously, unlike its predecessor the A10. The A11 also integrates an Apple-designed three-core graphics processing unit (GPU) with 30% faster graphics performance than the A10. Embedded in the A11 is the M11 motion coprocessor. The A11 includes a new image processor which supports computational photography functions such as lighting estimation, wide color capture, and advanced pixel processing. [from Wikipedia]

Tech  
Insights



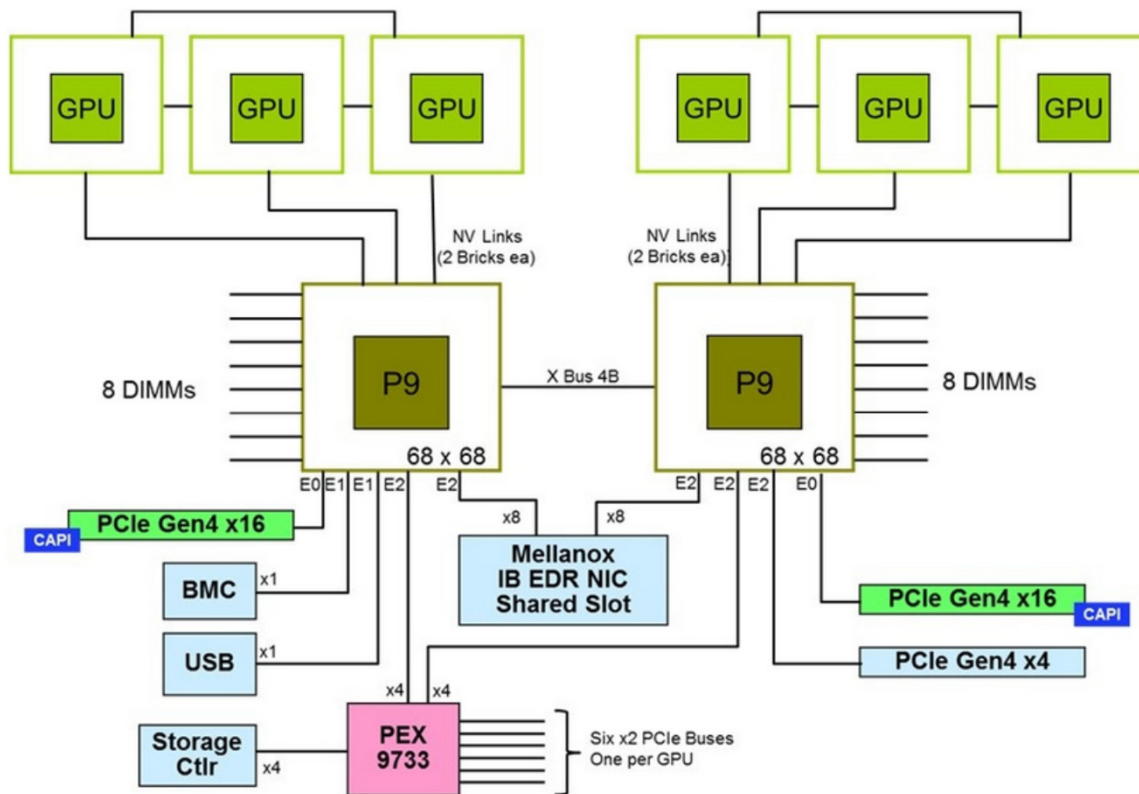
# Multicore Processor



IBM POWER9 Processor

*announced August 2016*

# POWER9 + GPU System





# Multicore Systems

---


Increased throughput: two CPUs can do more work in a given time than one CPU can!

Lower cost: multicore systems can share peripherals

More reliable: if one core fails, then the system can still function

Symmetric (SMP) vs. asymmetric (non-SMP) processing

- does every processor do the same thing for a particular program?
- or do different processors do different tasks?



# Computer Systems Architecture

---

Nowadays: multiprocessor (or multicore) systems are the norm

- note that “core” and “processor” are not synonymous
- multicore systems are truly the norm now



# Specialized Processors

---

GPU: graphics processing unit

- highly parallel hardware engine
- very good at solving a simple task efficiently
- more recently: much faster at solving the computations required for neural networks

TPU: Tensor processing unit

- custom hardware engine built by Google for TensorFlow
- to run machine-learning applications



# Operating System Structure

---

Most important characteristic: must support *multiprogramming*

- we don't want the CPU to be idle when a job is performing I/O
- in a multiprogrammed system, the OS will keep a set of jobs in a job pool (let's say this is in main memory)
- when it's turn for job A to run, the CPU switches its execution to A, and A starts to run
- when that job makes a request for I/O service, the OS switches to a different task – task B
- and so on
- basic goal: prevent the CPU from being idle (refer to the picture on Slide #11)



# Timesharing

---

Natural extension of multiprogramming is *timesharing*: enable many users to share system resources (most notably, the CPU)

- system switches rapidly between different users, giving them each a short time slice
- this gives all the users the illusion that they each have a dedicated machine!
- this will work effectively only if the response time for each user is reasonable and consistent
- this requires the OS to perform complex management of resources and requests
- *process*: a job that is loaded into memory and is ready for execution



# Memory Management

---

## physical memory

- the actual memory (the bits and bytes) on the machine

## virtual memory

- the view of memory that the OS uses and which a process sees
- can be much larger than physical memory
- the OS translates virtual-memory references to physical-memory references





# OS Operation

---

Simplest description of what the OS does:

- sit and wait for something to happen
- when something happens, take the appropriate action

Slightly more sophisticated description:

- let a process run until it makes an I/O request, or until timeout occurs, or until an unexpected condition in the software occurs\*
- then, give the CPU to a different process

\* what might this be? what's an unexpected condition that could occur in your software?



# Dual-Mode Operation

---

Multiprogrammed system: many different jobs and users, all using the same resources

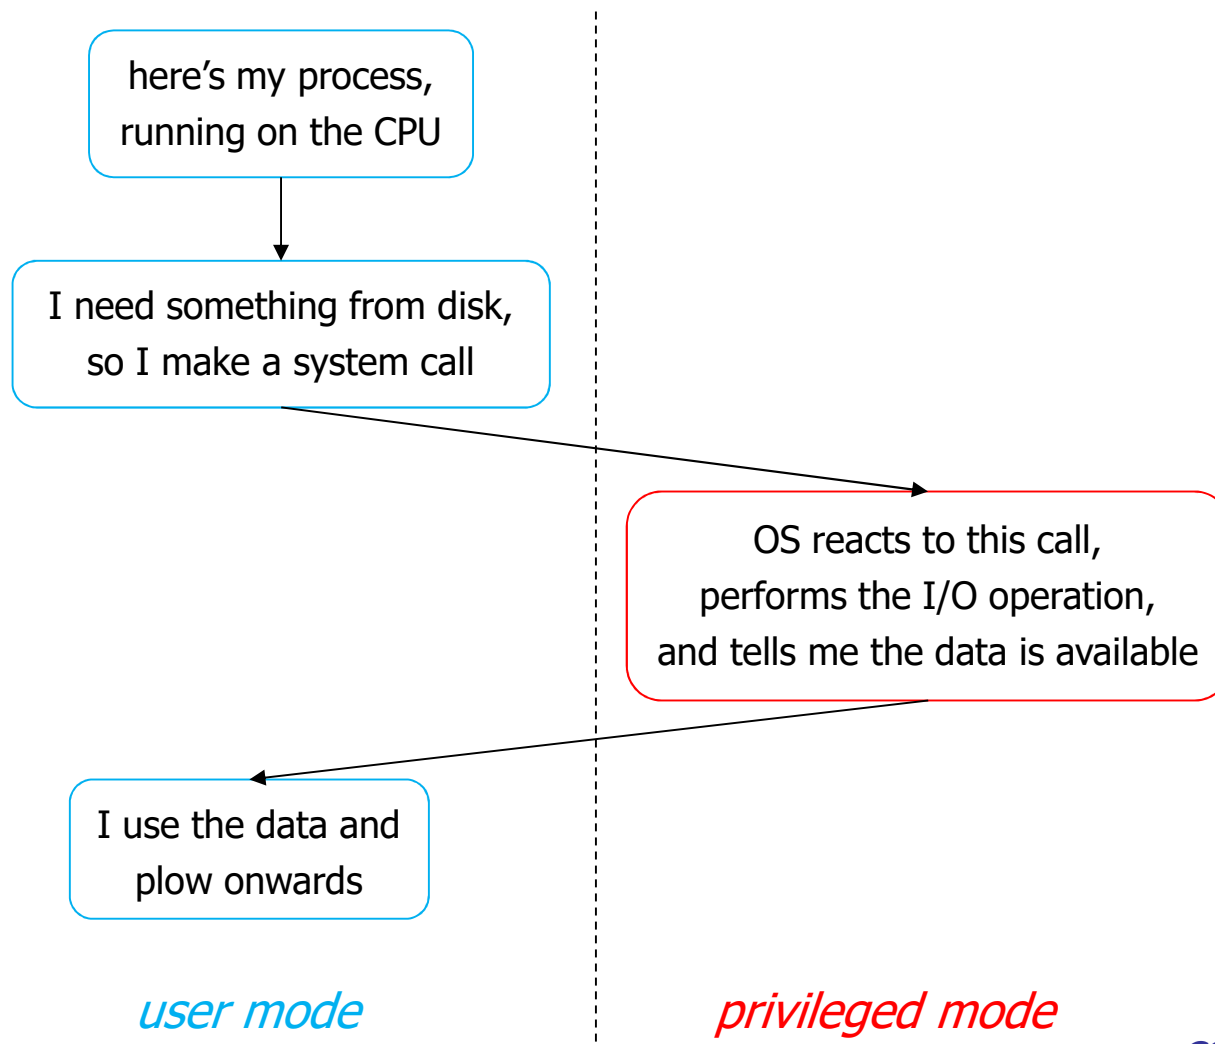
- need to prevent user A from accessing user B's memory!
- need to prevent user A from reading user B's private data!

This leads to the concept of dual-mode operation

- the OS needs special powers to facilitate this protection
- there are two modes of operation: user mode and kernel mode (aka privileged mode, system mode, supervisor mode)
- the mode is enforced in the hardware itself through a mode bit





# Dual-Mode Operation



## Memory: Extremely Simplified View of Memory Management

Think of the physical memory in the system

- the physical memory is the collection of memory boards in the system
- there's a finite amount of memory, and so the OS has to divvy the memory up among all the current users (actually, the current processes) on the machine





# Protection

---

If a process attempts an operation for which it does not have permission, the hardware intercedes!

- if I try to read memory that is not in my address space, I get a hardware exception

```
#include <stdio.h>

int main(int argc, char** argv) {
    int *ip;
    int dummy, idx;

    printf("hello!\n");

    for (idx = 0; idx < 1024; ++idx) {
        printf("idx = %d\n", idx);
        dummy = ip[1024 * idx];
    }
}
```



# Protection

---

Compiling and running this program:

```
squa11|/users/j/h/jhibbele/CS201
> gcc -g sigsegv.c
squa11|/users/j/h/jhibbele/CS201
> ./a.out
hello!
idx = 0
idx = 1
idx = 2
idx = 3
Segmentation fault
squa11|/users/j/h/jhibbele/CS201
>
```

What happens if I do  
this in Python?

What happens if I do  
this in Java?



# The Timer

---

OS must prevent a user process from keeping the CPU too long (when could this happen)?

This OS maintains a timer:

- when a user job gets the CPU, set the timer to  $t$  units
- when the timer expires, interrupt the job (if it still has the CPU)
- and at this point, a different job will be able to use the CPU

Question: who is permitted to modify the timer?



# A Process


---

Program: sequence of operations

Process: a program in execution


- process needs resources (CPU, I/O, memory), and it's the job of the OS to assign these resources to the process
- CPU executes one instruction at a time for a process
- when the process terminates, the OS cleans up after the process
- some processes belong to the OS itself
- a process can have several threads of execution (a thread is kind of like a piece of a process)





# Snapshot of Processes

```
root 1 0 0 Jan02 ? 00:01:00 /sbin/init
root 2 0 0 Jan02 ? 00:00:00 [kthreadd]
root 3 2 0 Jan02 ? 00:00:00 [klogd]
root 4 2 0 Jan02 ? 00:00:49 [ksftd]
root 5 2 0 Jan02 ? 00:00:00 [stopper]
root 6 2 0 Jan02 ? 00:00:20 [watchdog]
root 7 2 0 Jan02 ? 01:13:13 [events]
root 8 2 0 Jan02 ? 00:00:00 [events]
root 9 2 0 Jan02 ? 00:00:00 [events_long]
root 10 2 0 Jan02 ? 00:00:00 [events_poweroff]
root 11 2 0 Jan02 ? 00:00:00 [cgroup]
root 12 2 0 Jan02 ? 00:00:00 [khelper]
root 13 2 0 Jan02 ? 00:00:00 [netns]
root 14 2 0 Jan02 ? 00:00:00 [asyncmgr]
root 15 2 0 Jan02 ? 00:00:00 [jms]
root 16 2 0 Jan02 ? 00:00:57 [sync_supers]
root 17 2 0 Jan02 ? 00:00:02 [del-default]
root 18 2 0 Jan02 ? 00:00:00 [kintegrityd]
root 19 2 0 Jan02 ? 00:02:33 [kblockd]
root 20 2 0 Jan02 ? 00:00:00 [kacpid]
root 21 2 0 Jan02 ? 00:00:00 [kacpi_notify]
root 22 2 0 Jan02 ? 00:00:00 [kacpi_hotplug]
root 23 2 0 Jan02 ? 00:00:00 [ata_aux]
root 24 2 0 Jan02 ? 00:00:00 [ata_sff]
root 25 2 0 Jan02 ? 00:00:00 [ksuspend_usb]
root 26 2 0 Jan02 ? 00:00:00 [khubd]
root 27 2 0 Jan02 ? 00:00:00 [kserted]
root 28 2 0 Jan02 ? 00:00:00 [md]
root 29 2 0 Jan02 ? 00:00:00 [md_misc]
root 30 2 0 Jan02 ? 00:00:00 [l1mwatch]
root 33 2 0 Jan02 ? 00:00:04 [khungtaskd]
root 34 2 0 Jan02 ? 00:00:16 [kswapd0]
root 35 2 0 Jan02 ? 00:00:00 [kswapd]
root 36 2 0 Jan02 ? 00:01:56 [khugepaged]
root 37 2 0 Jan02 ? 00:00:00 [kio]
root 38 2 0 Jan02 ? 00:00:00 [crypto]
root 45 2 0 Jan02 ? 00:00:00 [kchrtod]
root 46 2 0 Jan02 ? 00:00:00 [jclibp]
root 48 2 0 Jan02 ? 00:00:00 [kpsmouse]
root 49 2 0 Jan02 ? 00:00:00 [usbhid_resumer]
root 50 2 0 Jan02 ? 00:00:00 [deferwq]
root 82 2 0 Jan02 ? 00:00:00 [kdmrmov]
root 83 2 0 Jan02 ? 00:00:00 [kstrp]
root 112 2 0 Jan02 ? 00:00:00 [tts_swap]
root 209 2 0 Jan02 ? 00:00:00 [scsi_ah_0]
root 210 2 0 Jan02 ? 00:00:00 [scsi_ah_1]
root 216 2 0 Jan02 ? 00:00:00 [scsi_ah_2]
root 217 2 0 Jan02 ? 00:00:00 [vmm_pvcsl_wq_2]
root 341 2 0 Jan02 ? 00:00:00 [kdmflush]
root 343 2 0 Jan02 ? 00:00:00 [kdmflush]
root 361 2 0 Jan02 ? 00:04:59 [jbd2/de-0-8]
root 362 2 0 Jan02 ? 00:00:00 [ext4-dio-unwrt]
root 437 1 0 Jan02 ? 00:00:18 /sbin/udev -d
root 577 2 0 Jan02 ? 00:04:28 [vmmemc1]
root 748 2 0 Jan02 ? 00:00:00 [jbd2/sd1-8]
root 749 2 0 Jan02 ? 00:00:00 [ext4-dio-unwrt]
root 838 2 0 Jan02 ? 00:01:02 [kaudit]
root 1358 2 0 Jan02 ? 00:04:27 [flush-733]
root 1418 1 0 Jan02 ? 00:05:07 auditd
root 1452 1 0 Jan02 ? 00:34:28 /usr/sbin/rsyslogd
root 1468 1 0 Jan02 ? 00:07:09 /usr/sbin/rsyslogd -i /var/run/syslogd.pid -c 5
dbus 1539 1 0 Jan02 ? 00:00:00 dbus-daemon --system
root 1591 2 0 Jan02 ? 00:04:12 [rpciod]
root 1593 2 0 Jan02 ? 00:00:00 [kslowd000]
root 1594 2 0 Jan02 ? 00:00:00 [kslowd001]
root 1595 2 0 Jan02 ? 00:00:12 [nfsiod]
root 1596 2 0 Jan02 ? 00:00:00 [rfsv4.0-svc]
root 1738 1 0 Jan02 ? 00:05:29 sendmail: accepting connections
root 1747 1 0 Jan02 ? 00:00:01 sendmail: Queue runner001:00:00 for /var/spool/clientmqueue
root 1767 1 0 Jan02 ? 00:00:00 /usr/bin/svmon --daemon --pid-file=/var/run/svmon.pid -d -r /usr/local/subversion
root 1779 1 0 Jan02 ? 00:03:34 cron
root 1836 1 0 Jan02 ? 00:00:00 /usr/sbin/atd
root 1849 1 0 Jan02 tty1 00:00:00 /sbin/mingetty /dev/tty1
root 1851 1 0 Jan02 tty2 00:00:00 /sbin/mingetty /dev/tty2
root 1853 1 0 Jan02 tty3 00:00:00 /sbin/mingetty /dev/tty3
root 1855 1 0 Jan02 tty4 00:00:00 /sbin/mingetty /dev/tty4
root 1859 1 0 Jan02 tty5 00:00:00 /sbin/mingetty /dev/tty5
root 1861 1 0 Jan02 tty6 00:00:00 /sbin/mingetty /dev/tty6
root 3036 1 2 11:43 ? 00:00:07 /usr/bin/python /usr/bin/fail2ban-server -b -s /var/run/fail2ban/fail2ban.sock -p /var/run/fail2ban/fail2ban.pid -x
jhibbele 4072 27208 0 11:48 pts/1 00:00:00 cupsd -C /etc/cups/cupsd.conf
root 7636 1 0 Jan08 ? 00:03:51 /usr/sbin/sshd
root 7670 1 0 Jan08 ? 00:02:17 automount --pid-file /var/run/automounts.pid
rpc 7805 1 0 Jan08 ? 00:00:22 rpcbind -w
root 9029 1 0 Jan29 ? 02:22:48 /usr/sbin/vmtoolsd
ntp 11155 1 0 Jan29 ? 00:00:26 ntpd -u ntp:ntp -g /var/run/ntpd.pid -g
root 13162 7670 0 Jun11 ? 00:00:00 sshd: xwu [priv]
xwu 13195 13162 0 Jun21 ? 00:00:00 sshd: xwu@pts/0
xwu 13196 13195 0 Jun21 pts/0 00:00:00 tcsh
root 24290 1 0 Jun12 ? 00:18:34 /var/cfengine/bin/cf-execd
root 24588 1 0 Jun12 ? 00:06:26 /var/cfengine/bin/cf-monitord
root 24702 1 0 Jun12 ? 00:28:43 /var/cfengine/bin/cf-serverd
root 27195 7670 0 Jun16 ? 00:00:00 sshd: jhibbele [priv]
jhibbele 27207 27195 0 Jun17 ? 00:00:00 sshd: jhibbele@pts/1
jhibbele 27208 27207 0 Jun17 pts/1 00:00:00 -bash
rpcuser 28929 1 0 Jan08 ? 00:00:00 rpc.statd
root 32645 2 0 Jun07 ? 00:00:00 [flush-2-24]
root 32748 437 0 Jan29 ? 00:00:14 /sbin/udev -d
root 32749 437 0 Jan29 ? 00:00:00 /sbin/udev -d
```



# Processes: What the OS Does

---

1. schedule processes and threads on the CPUs
2. create and delete user and system processes
3. suspend and resume processes
4. provide mechanisms for process synchronization
5. provide mechanisms for process communication



# Memory Management

---

Remember: bit vs. byte. KB, MB, GB.

Main memory: the hardware memory that the CPU can access directly

In order to access data on disk, the data must first be brought into main memory

- note that this includes the program instructions themselves!




# Memory Management

---

As the program executes, the memory references in the program must be mapped to real physical memory

The OS will keep several programs (and their data) in memory simultaneously

- and manage the memory and the processes



# Memory: What the OS Does

---

1. keep track of which parts of memory are being used, and who is using them
2. decide which processes and what data to move into and out of memory
3. allocate and deallocate memory as needed



# Storage Management

---

File: collection of related information, in some particular format

A file lives on some storage medium in the system

- let's say on a hard disk drive

The hard disk drive itself has some particular organization to it that the OS imposes

The OS additionally imposes access control on files



# Storage: What the OS Does

---

1. create and delete files
2. create and delete directories to organize files
3. support primitive operations for manipulating files and directories
4. map files onto secondary storage (i.e. disk drive)
5. back up files on stable storage media



# Mass-Storage Management

---

Disks and beyond; operations are:

1. managing free space
2. allocating storage
3. scheduling disk accesses





# Caching

---

Just a quick mention—idea is to keep the data that you need close to you (see Slide #11)

Hardware memory cache, managed by the hardware

Main memory is a cache for secondary memory; managed by the OS (secondary memory—disk drive or flash memory)



# I/O Subsystem

---

The I/O subsystem abstracts away the details of the I/O devices themselves

Key activities of the I/O subsystem of the OS:

1. buffer, cache, and spool data from the I/O devices to/from main memory
2. provide a general device-driver interface
3. provide drivers for specific hardware devices



# Protection and Security

---

Protection: mechanism for controlling access of processes (and users) to system resources

Security: prevention of unauthorized access to system resources

The system assigns each user a unique ID

- and then grants certain privileges to each ID
- additionally, the system can form groups of users with certain privileges associated with the group



# Virtualization

---

Virtualization is the enablement in software of a “virtual machine”

- a self-contained environment on a host operating system that appears to be a separate machine, possibly running a different OS
- useful for running programs written for one OS on a machine that uses a different OS
- also provides the basic infrastructure for cloud computing



# Distributed Systems

---

## Distributed system

- a collection of physically distinct computer systems that are networked together
- provide users access to shared resources
- might provided a distributed file system (e.g., NFS or AFS)



# Kernel Data Structures

---

Review: here are some key data structures that an OS uses


1. lists, stacks, queues
2. trees
3. hashmaps (dictionaries)
4. bitmaps

# Computing Environments



What company is this?

# Computing Environments



What company is this?



# Computing Environments



What company is this?

# Computing Environments



What company is this?



# Computing Environments

---

In the old days: everyone had a “personal computer”

- connected to a network
- all the big work was done on big batch servers

Now we have laptops

And mobile systems

- tablet or “smart phone”
- does a smart phone have an OS? How about a Kindle?

Connection models

- Client/server, peer-to-peer



# Cloud Computing

---

Public cloud, private cloud

Kinds of cloud computing

- software as a service (SaaS)
- platform as a service (PaaS)
- infrastructure as a service (IaaS)



# Real-Time Embedded Systems

---

Usually created to accomplish very specific tasks

- might be a primitive OS
- or might be a full OS, e.g., a specialized version of Linux

Often, the #1 performance criterion is response time

- think of a robot arm in a factory
- or a self-driving car (!!!)
- for non-real-time systems, the variance of the response time is important
- for real-time systems, the maximum response time is key



# Finally, Open-Source Software

---

Linux is an open-source system

- and there are many flavors of Linux

Darwin, the kernel of macOS, is open source

Android is essentially open source