



# *Threads & Concurrency*

---

## CS201 Lecture 4

Jason Hibbeler  
University of Vermont  
Spring 2019

### Bohemian Hippie Threads

1.93k Pins · 4.56k Followers

Hippie at heart!! Collection of boho, gypsy, hippie styles.



by Gypsy Dreams

Bohemian style

Woman fashion

Bohemian fashion



# Thread

---

Basic unit of CPU utilization

- an ID
- a program counter
- register values
- a stack

It's like the essential essence of a process

- remember that the process has data, text (code), system resources (such as files)
- all threads in a process share these resources



# Thread

---

General statement:

*If a process has more than one thread, and these threads can execute simultaneously\*, then the process should be quicker*

\*we'll discuss what this actually means



# Threads in Action

---

Think about a word-processing program

- at the same time you're typing text, the program is doing spellchecking for you
- it might also do an automatic save while you're typing text

Or an even better example: an IDE

- with all of the suggestions, checking, and cross-referencing that an IDE does under the covers while you are typing in program text



# Web Server

---

Another example: a web server

## Structure of web server

- listen for a request
- when a request appears, run code to handle that request

If the server has a single thread of execution, then it won't be able to handle new requests while it's tending to an existing request

- why not have the server just fork off a new process to handle each new request? We will answer this next.



# Benefits of Multithreading

---

1. Responsiveness
2. Resource sharing
3. Economy
4. Scalability



# Responsiveness

---

Without threading, then if an application blocks or is performing a lengthy computation, then the whole program blocks

- UI must always be responsive
- put the computation in a separate thread
- put the calls that have unknown response time in a separate thread (think about any application that needs to get input from the web)



# Resource Sharing

---

Resource sharing among processes is somewhat complex for application developers (shared memory, message passing)

- threads share memory and system resources
- this makes programming multithreaded applications somewhat less awkward than multi-process applications (at least from the standpoint of sharing resources)





# Economy

---

Since threads inherit the system resources of the parent process, it's faster and cheaper to create threads than it is to create new processes

- and the context switch between threads is faster as well



# Scalability

---


In a multicore machine, different threads of a process can run on different cores

- whereas a single-threaded process can run only on a single core



# Concurrency vs. Parallelism

Think about how a single-core system processes tasks (threads) for a user. Here is an idealized diagram showing the processing of six threads in the system:



What's the throughput per unit time?

- it's  $1 / \text{numthreads}$


The threads are being executed *concurrently*

- all of them are making progress



# Concurrency vs. Parallelism

Now suppose we have two CPUs (cores):



What's the throughput per unit time? It's 2 / numthreads

- twice as fast!

Two threads are being executed *in parallel*

- if we can swing this, it's clearly better

*In what circumstances can we actually do this?*



# Multithreading: Challenges

---

Creating multithreaded applications is challenging.

A general observation:

*It is much more difficult to rewrite a single-threaded application to be multithreaded than it is to design the application to be multithreaded from the beginning.*



# Multithreading: Challenges

---

Identifying tasks:

- in the processing of the application, we want to have separate tasks that can execute independently of each other.


Is this easy to accomplish?

- it depends on the application

# Multithreading: Challenges

Balance: even if we can decompose the application into separate, independent tasks, we want to insure that the tasks all perform approximately the same amount of work.

More specifically: the work they each perform should take approximately the same amount of time.





- example: suppose we are doing some processing on a triangular matrix
- here, the yellow tasks have only half as much work as blue tasks

# Multithreading: Challenges

Ten threads start together, each on its own core

- but four of the threads have only half as much work



- if each CPU is mapped to a core, then four of the cores are idle for half of the time
- this gives us a total utilization of 80%





# Multithreading: Challenges

---

## Data splitting:

- in order to enable separate tasks to run independently, we need to be able to partition the data so that the tasks can access\* the data that they need in an independent fashion

\*key point here: if the tasks merely read the data, then we don't have to do anything special; but if tasks modify the data then we have to arrange things so that the data can be independent



# Problem Decomposition

---

Example: adding two vectors

- worker thread #1 can add  $A[0..999] + B[0...999]$
- worker thread #2 can add  $A[1000..1999] + B[1000...1999]$
- etc.

Each thread can do its work independently of the other threads



# Problem Decomposition

---

Another example: finding the maximum of the values in an array

- worker thread #1 can find the max in  $A[0..999]$
- worker thread #2 can find the max in  $A[1000..1999]$
- etc.

But what has to happen next?



# Problem Decomposition

---

Another example: finding the maximum of the values in an array

- worker thread #1 can find the max in  $A[0..999]$
- worker thread #2 can find the max in  $A[1000..1999]$
- etc.

But what has to happen next?

- another thread has to look at the max value that each thread has found and take the max of that
- so this “master thread” can do its work only after the “worker threads” have finished



# Problem Decomposition

---

Another example: sorting an array

- worker thread #1 can sort the values in  $A[0..999]$
- worker thread #2 can sort the values in  $A[1000..1999]$
- etc.

What has to happen next?



# Problem Decomposition

---

Another example: sorting an array

- worker thread #1 can sort the values in  $A[0..999]$
- worker thread #2 can sort the values in  $A[1000..1999]$
- etc.

What has to happen next?

- the master thread has to merge all of the sorted subarrays together



# Multithreading: Challenges

---

## Thread-local storage

- Threads inherit the data space of the process that created them
- But it's also important for threads to have their own private storage (we'll see examples)
- And the casual modification of a non-thread-local variable by a thread can cause unpredictable behavior in other threads that are using that variable


# Multithreading: Challenges

## Data dependency:

- if task  $T_k$  depends on data that is produced by task  $T_j$ , then  $T_k$  can execute only after  $T_j$  is complete

## Furthermore:

- we must impose a synchronization regime so that  $T_k$  knows when  $T_j$  has produced the data that  $T_k$  needs.







# Multithreading: Challenges

---

## Testing and debugging:

- instead of just a single thread of execution to trace, we will have several threads
- and the order in which they execute could be indeterminate (or, more importantly, unpredictable)

## Summary:

- does it seem like it would be difficult to design and develop multithreaded application programs?
- yes!



# Two Basic Models of Parallelism

---

## **Data parallelism**

- split up the data across tasks; create a thread for each task
- each task is identical but operates on its own subset of the data.

Simple example is the addition of two vectors

- task #1 sums  $A[0:999] + B[0:999]$
- task #2 sums  $A[1000:1999] + B[1000:1999]$
- etc.
- so each thread is working safely in its own region of A and B



# Two Basic Models of Parallelism

---

## **Task parallelism**

- the two tasks are different and perform different operations.
- simple example: one task computes the mean of the values in an array
- and a second task computes the max of the values



# Models of Multithreading

---

**user thread:** the user application manages the control of the thread

- here, your code actually creates the threads

**kernel thread:** the kernel manages the control of the thread

- the OS itself creates the threads



# Models of Multithreading

---

The OS schedules kernel threads



- in other words, it's the kernel threads that get put onto a CPU
- the threading library (the part of the OS that actually provides the threading interface for application developers) decides how to map user threads to kernel threads

# Models of Multithreading

## many-to-one vs. one-to-one

- the OS is configured to do one of these

*K is kernel thread  
U is user thread*




Consequence of this mapping:

- many-to-one cannot actually achieve parallelism

# Models of Multithreading

Most effective model is many-to-many

- so if one user thread blocks, the OS can schedule a different kernel thread for execution





# Thread Libraries

---

A thread library provides the APIs for the application developer to create and manage threads

- we'll look at a couple of python examples (on gitlab)
- we'll also look at a Unix pthread example: given an integer  $N$ , sum all the integers from 1 to  $N$





# Unix Example: pthreads

```
#include <pthread.h>
#include <stdio.h>

#define NUMCHILDREN 2

void *runner(void *param);

typedef struct {
    int lowVal;
    int highVal;
    int sum;
} SumStruct;

int main(int argc, char *argv[]) {
    SumStruct data[NUMCHILDREN]; // holds data we want to
                                // give to child thread
    pthread_t tid[NUMCHILDREN]; // thread identifier
    pthread_attr_t attr; // thread attributes
    int bigSum;
    int i;

    data[0].lowVal = 1;
    data[0].highVal = 50;
    data[1].lowVal = 51;
    data[1].highVal = 100;

    // get default thread attributes
    pthread_attr_init(&attr);
```

```
    for (i=0; i<NUMCHILDREN; ++i) {
        // create child thread
        pthread_create(&tid[i], &attr, runner,
                      &data[i]);
    }

    for (i=0; i<NUMCHILDREN; ++i) {
        // wait for the child threads to terminate
        pthread_join(tid[i], NULL);
    }

    bigSum = 0;
    for (i=0; i<NUMCHILDREN; ++i) {
        bigSum = bigSum + data[i].sum;
    }

    printf("sum is %d\n", bigsum);
    return(0);
}
```

this code is in gitlab, in pthreads-example.c,  
under Examples/



# Unix Example: pthreads

---

```
void *runner(void *param) {
    SumStruct *data;
    int i, sum;

    sum = 0;
    data = (SumStruct *) param;

    printf("(R) I am runner; will sum integers from %d to %d\n",
           data->lowVal, data->highVal);

    for (i=data->lowVal; i<=data->highVal; ++i)
        sum = sum + i;
    data->sum = sum;

    printf("(R) sum is %d\n", data->sum);

    pthread_exit(0);
}
```

this code is in gitlab, in pthreads-example.c,  
under Examples/



# Threads in Java

---

How is multithreading actually supported in Java?

Java code executes in a self-contained environment called the Java Virtual Machine.

So, this means that the JVM is managing the threads for us

See Java example in textbook (section 4.4.3).



# Threads in Python

---

Python has a structure called the “global interpreter lock” (GIL)

- this prevents more than one thread from executing at a single time
- so naïve threading in Python won't provide any speed-up
- if you read a discussion of threading in Python, you'll see that knowledgeable people say “don't do threading in Python”
- but it's a good way to prototype an algorithm or a program
- and there is a different way to achieve parallelism in Python



# Skip

---

We will skip Section 4.5 and 4.6 and 4.7.